


Amazon SageMaker

Train, Tune and Deploy



Deep Dive

In this session: Train, Tune and Deploy models

Prepar
e

Build

Train & Tune

Deploy &
Manage

Web-based IDE for ML

Automatically build and train models

Fully managed data processing jobs and data labeling workflows

`10101101
0
01010101
0
00001111
0`

One-click collaborative notebooks and built-in, high performance algorithms and models



One-click training



Debugging and optimization



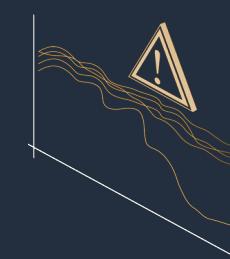
Visually track and compare experiments



One-click deployment and auto-scaling




Automatically spot concept drift



Add human review of predictions



Fully managed with auto-scaling for 75% less



Collect and prepare training data

Choose or build an ML algorithm

Set up and manage environments for training

Train, debug, and tune models

Manage training runs

Deploy model in production

Monitor models

Validate predictions


Scale & manage the production environment




Train ML Model in SageMaker




Learning Theory Fundamentals



Overfitting




Underfitting



Underfitting

Underfitting: The model is not good enough to describe the relationship between the input data (x) and output (y).



The model is too simple to capture the input/output relationship.

It will have poor training and test performance.

Overfitting

Overfitting: Model memorizes or imitates training data and doesn't generalize well with new “unseen” data (test data).




Having too complex models for simple problems can cause overfitting.

The model picks up the noise instead of the underlying relationship.

We will see good scores in training data but poor in test data.

Appropriate Fitting

Appropriate fitting: It captures the general relationship between the input data (x) and output (y).



Amazon SageMaker Algorithms



- Matrix factorization
- Regression
- Principal component analysis
- K-means clustering
- Gradient boosted trees
- And more!

17 Built-in algorithms



Bring your own script
(Amazon SageMaker managed container)




Bring your own
algorithm
(you build the
Docker container)



Subscribe to
Algorithms and
Model Packages
on AWS
Marketplace



Training Jobs




Every model run on a SageMaker training job has its own **ephemeral cluster**.

That means you have a dedicated EC2 instance alive for the **number of seconds** your model needs to train.

This **cluster comes down immediately** after the model finished training.




Amazon SageMaker training service




Amazon SageMaker | Training

Programming Model



End-to-End Flow



Amazon SageMaker | Training

Use built-in algorithms


```
linear = Estimator('linear-learner',
    train_instance_count=1,
    train_instance_type='ml.c4.xlarge',
    output_path=output_location,
    sagemaker_session=sess)
```

```
linear.set_hyperparameters(
    feature_dim=784,
    predictor_type='binary_classifier',
    mini_batch_size=200)
```

```
linear.fit({'train': s3_train_data})
```

Images

- Xgboost
- PCA
- DeepAR
- BlazingText
- Image classification
- ...
- Object Detection



```
In [ ]: bt_model = sagemaker.estimator.Estimator(container,
                                                role,
                                                train_instance_count=1,
                                                train_instance_type='ml.c4.4xlarge',
                                                train_volume_size = 30,
                                                train_max_run = 360000,
                                                input_mode= 'File',
                                                output_path=s3_output_location,
                                                sagemaker_session=sess)
```

Number of EC2 instances

Type of EC2 instances

Disk space



Algorithm Container

```
In [ ]: bt_model = sagemaker.estimator.Estimator(container,
                                                role,
                                                train_instance_count=1,
                                                train_instance_type='ml.c4.4xlarge',
                                                train_volume_size = 30,
                                                train_max_run = 360000,
                                                input_mode= 'File',
                                                output_path=s3_output_location,
                                                sagemaker_session=sess)
```

SageMaker Estimator

Execution Role

Cluster comes online
Logs to CloudWatch
Monitor via console or notification stream



Confusion Matrix

Your Model's Predictions

Your Labeled Data

| | | Your Model's Predictions | |
|-------------------|----------|--------------------------|----------------|
| | | Positive | Negative |
| Your Labeled Data | Positive | True Positive | False Negative |
| | Negative | False Positive | True Negative |

Recall (Yellow circle, green checkmark)

Precision (Yellow circle, red X)

Evaluating classification models

| | | Predicted Response | |
|---------------|---------|--------------------|----------------|
| | | $\hat{y} = 1$ | $\hat{y} = 0$ |
| True Response | $y = 1$ | True Positive | False Negative |
| | $y = 0$ | False Positive | True Negative |

Recall (sensitivity):

Precision:

Precision: Accuracy of a predicted positive outcome.

F1 - Score: Harmonic mean of precision and recall.

Recall (sensitivity): Measures the strength of the model to predict a positive (1) outcome.

(good)



Tune ML Model in SageMaker





Amazon SageMaker Automatic Model Tuning

Hyperparameter Optimizer



Decision Trees

Tree depth
Max leaf nodes
Gamma
Eta
Lambda
Alpha
...


Neural Networks

Number of layers
Hidden layer width
Learning rate
Embedding dimensions
Dropout
...


“Hyperparameters”

(algorithm parameters that significantly affect model quality)

Hyperparameter Tuning Jobs



Hyperparameter Tuning Jobs



Different Hyperparameter Settings


Round Count

2/7



Hyperparameter Tuning Jobs

Round 7/7



How do I set up a hyperparameter tuning job?



Objective Metric

δ

Hyperparameter
s & Ranges



Job Specs



Can I use hyperparameter tuning with my own model?

Yes!!!



1

Built-in
Algorithms



Docker



Script Mode

Fully Customizable



Can I use hyperparameter tuning with my own model?

Setting the hyperparameters

```
In [5]: hyperparameters = dict(batch_size=32, data_augmentation=True, learning_rate=.0001,  
                           width_shift_range=.1, height_shift_range=.1, epochs=1)  
hyperparameters
```

```
Out[5]: {'batch_size': 32,  
         'data_augmentation': True,  
         'learning_rate': 0.0001,  
         'width_shift_range': 0.1,  
         'height_shift_range': 0.1,  
         'epochs': 1}
```

Docker



Script Mode



1. Pick hyperparameters and ranges

```
: hyperparameter_ranges = {'eta': ContinuousParameter(0, 1),  
                           'min_child_weight': ContinuousParameter(1, 10),  
                           'alpha': ContinuousParameter(0, 2),  
                           'max_depth': IntegerParameter(1, 10)}
```

2. Pick objective metric

```
: objective_metric_name = 'validation:auc'
```

3. Pick job parameters

```
: tuner = HyperparameterTuner(xgb,  
                               objective_metric_name,  
                               hyperparameter_ranges,  
                               max_jobs=20,  
                               max_parallel_jobs=3)
```



What if I need all my jobs tuned at the same time?

Random search



What if I need all my jobs tuned at the same time?

Bayesian Search



Random Search



What if I need all my jobs tuned at the same time?

Random search

```
{
    "ParameterRanges": {...},
    "Strategy": "Random",
    "HyperParameterTuningJobObjective": {...}
}
```

```
tuner = HyperparameterTuner(
    sagemaker_estimator,
    objective_metric_name,
    hyperparameter_ranges,
    max_jobs=20,
    max_parallel_jobs=20,
    strategy="Random"
)
```



Amazon SageMaker Automatic Model Tuning


```
from sagemaker.tuner import IntegerParameter, CategoricalParameter, ContinuousParameter, Hyperparameter
```

```
hyperparameter_ranges = {  
    'batch_size': CategoricalParameter([16, 32, 64, 128]),  
    'learning_rate': ContinuousParameter(0.0009, 0.  
    'epochs': IntegerParameter(10, 100)}
```

```
objective_metric_name = 'Validation-accuracy'  
metric_definitions = [{Name: 'Validation-accuracy',  
    'Regex': 'Validation-accuracy=([0-9\\.]+)' }]
```

```
tuner = HyperparameterTuner(estimator,  
    objective_metric_name,  
    hyperparameter_ranges,  
    metric_definitions,  
    max_jobs=20,  
    max_parallel_jobs=4)
```

```
tuner.fit(dataset_location, job_name=job_name)
```



Pro Tips


- S3 bucket must be in the same region as your training job
- Use SageMaker Search to compare results of previous jobs
- Run training jobs in parallel from your notebook
- Soft limit on number of instances used for training job
- Check docs on which hyperparameters are valid for tuning
- When bringing your own algorithm, set hyperparameters as JSON object




Model Deployment and Inferencing with SageMaker




Amazon SageMaker Deployment Options



Amazon SageMaker hosting service



SageMaker Endpoints



Deployment / Hosting
Amazon SageMaker ML
Compute Instances



Command Line - Creating and Scaling Model Endpoints



Easy deployment to production REST API



Scalable, high throughput, and high reliability



Creating endpoints

Model

```
aws sagemaker create-model  
  --model-name modell  
  --primary-container '{"Image": "123.dkr.ecr.amazonaws.com/algo",  
                      "ModelDataUrl": "s3://bkt/modell.tar.gz"}'  
  --execution-role-arn arn:aws:iam::123:role/me
```

Endpoint configuration


```
aws sagemaker create-endpoint-config  
  --endpoint-config-name modell-config  
  --production-variants '{"InitialInstanceCount": 2,  
                        "InstanceType": "ml.m4.xlarge",  
                        "InitialVariantWeight": 1,  
                        "ModelName": "modell",  
                        "VariantName": "AllTraffic"}'
```

Endpoint

```
aws sagemaker create-endpoint  
  --endpoint-name my-endpoint  
  --endpoint-config-name modell-config
```



Updating endpoints



Blue-green
deployments mean
no scheduled
downtime



Deploy one or more
models behind the
same endpoint



Updating endpoints

New model

```
aws sagemaker create-model  
--model-name model2  
--primary-container '{"Image": "123.dkr.ecr.amazonaws.com/algo",  
                     "ModelDataUrl": "s3://bkt/model2.tar.gz"}'  
--execution-role-arn arn:aws:iam::123:role/me
```

New
endpoint
configuration


```
aws sagemaker create-endpoint-config  
--endpoint-config-name model2-config  
--production-variants '{"InitialInstanceCount": 2,  
                      "InstanceType": "ml.m4.xlarge",  
                      "InitialVariantWeight": 1,  
                      "ModelName": "model2",  
                      "VariantName": "AllTraffic"}'
```

Same
endpoint

```
aws sagemaker update-endpoint  
--endpoint-name my-endpoint  
--endpoint-config-name model2-config
```



Reduced risk deployments



Incrementally
retrain models with
new data



Try new models and
improved
algorithms



Reduced risk deployments

Two-model endpoint configuration

```
aws sagemaker create-endpoint-config  
  --endpoint-config-name both-models-config  
  --production-variants '[{"InitialInstanceCount": 2,  
    "InstanceType": "ml.m4.xlarge",  
    "InitialVariantWeight": 95,  
    "ModelName": "model1",  
    "VariantName": "model1-traffic"},  
   {"InitialInstanceCount": 2,  
    "InstanceType": "ml.m4.xlarge",  
    "InitialVariantWeight": 5,  
    "ModelName": "model2",  
    "VariantName": "model2-traffic"}]'
```

Same endpoint

```
aws sagemaker update-endpoint  
  --endpoint-name my-endpoint  
  --endpoint-config-name both-models-config
```

Swap

```
aws sagemaker update-endpoint-weights-and-capacities  
  --endpoint-name my-endpoint  
  --desired-weights-and-capacities '{"VariantName": "model1",  
    "DesiredWeight": 5}'
```



Automatic scaling endpoints

SageMaker console settings:

- Min and max instances
- Target invocations per instance
- Scaling cooldowns

Variant automatic scaling [Learn more](#)

| Variant name | Instance type | Current instance count | Current weight |
|--------------|---------------|------------------------|----------------|
| AllTraffic | ml.p2.xlarge | 2 | 1 |

Minimum instance count Maximum instance count
2 - 5

IAM role
Amazon SageMaker uses the following service-linked role for automatic scaling. [Learn more](#)
`AWSServiceRoleForApplicationAutoScaling_SageMakerEndpoint`

Built-in scaling policy [Learn more](#)


| Policy name | |
|--|--|
| SageMakerEndpointInvocationScalingPolicy | |

Target metric Target value
`SageMakerVariantInvocationsPerInstance` 800

Scale in cool down (seconds) - optional Scale out cool down (seconds) - optional
120 60



Scaling criteria



Algorithms have
different memory,
CPU, or GPU
requirements



Automatically scale
based on endpoint
instance's Amazon
CloudWatch metrics

Creating an automatic scaling policy

Variant


```
aws application-autoscaling register-scalable-target
  --service-namespace sagemaker
  --resource-id endpoint/my-endpoint/variant/model2
  --scalable-dimension sagemaker:variant:DesiredInstanceCount
  --min-capacity 2
  --max-capacity 5
```

Policy

```
aws application-autoscaling put-scaling-policy
  --policy-name model2-scaling
  --service-namespace sagemaker
  --resource-id endpoint/my-endpoint/variant/model2
  --scalable-dimension sagemaker:variant:DesiredInstanceCount
  --policy-type TargetTrackingScaling
  --target-tracking-scaling-policy-configuration
  '{"TargetValue": 50,
   "CustomizedMetricSpecification":
     {"MetricName": "CPUUtilization",
      "Namespace": "/aws/sagemaker/Endpoints",
      "Dimensions":
        [{"Name": "EndpointName", "Value": "my-endpoint"},
         {"Name": "VariantName", "Value": "model2"}],
      "Statistic": "Average",
      "Unit": "Percent"}}'
```



Scale by utilization





Customer


Endpoints are
easy to spin up,
but what if I need
to serve
predictions on a
schedule?



Cloud Architect

Batch
transform!

Batch Transform Common Design Pattern



SageMaker Batch Transform

Amazon SageMaker > Training jobs > knn-190612-2330-009-b61fb557

knn-190612-2330-009-b61fb557

Clone

Create model package

Stop

Create model

1.

Job settings

2.

Model settings

Model name

my-model-name

Maximum of 63 alphanumeric characters. Can include hyphens (-), but not spaces. Must be unique within your account in an AWS Region.

3.

Amazon SageMaker > Batch transform jobs > Create batch transform job

Create batch transform job

A transform job uses a model to transform data and stores the results at a specified location. [Learn more](#)

Batch transform job configuration

© 2020, Amazon

Job name

my-batch-transform-job

Maximum of 63 alphanumeric characters. Can include hyphens (-), but not spaces. Must be unique within your account in the same AWS Region.

Model name

my-model-name

Find model

Maximum of 63 alphanumeric characters. Can include hyphens (-), but not spaces. Must be unique within your account in the same AWS Region.

The cluster will spin down immediately after the job is finished.



Choosing Your Deployment Option

| | Scaling | Management | Flexibility | Cost | Latency |
|--------------------|-----------------------------------|-------------|----------------------|---|----------------------|
| SageMaker Endpoint | Auto-scaling enabled | AWS-managed | Depends on SM Docker | Higher: Cluster runs always and can scale based on demand | Lowest latency |
| Batch Transform | Configure size of cluster per run | AWS-managed | Depends on SM Docker | Lower: Cluster is decommissioned after use | 3-4 minute wait time |



Pro Tips

- Turn your endpoints off when not in use with Lambda
- Use inference pipelines for pre and post processing
- Bring more models in your own Docker container
- Consider the size of data hitting a single endpoint
- You can train your model elsewhere and host in SageMaker




Thank you!!!



Appendix




Gradient Descent



At each step, we take a step towards minimum by the value of gradient*learning rate.


Towards mid sections, gradient (slope) becomes smaller, therefore, steps become smaller.

Gradient Descent




Learning rate

If learning rate is small, it can converge slowly.



If learning rate is too large, gradient descent can overshoot the minimum.



Bagging

Bagging is an ensemble learning technique where we:

1. Draw bootstrap samples from the training set
2. Train a ML model on each sample

